Informatics II

Tutorial Session 13

Thursday, 27th of May 2021

Discussion of Exercise 12,
Graph Theory,
Preparation for Final Exam

09.00 – 11.45

Online on Zoom (Meeting ID: 970 3019 1915)

Agenda

- Intro, Expectations and Wishes, Administrivia
- Recitation of Graphs and Review of Exercise 12 (interleaved)
  - Graph Representation
  - Graph Traversals: DFS, BFS,
  - Topological Sorting
- Exam Preparation
- Summary and Wrap-Up

Introduction

- What shall we do today?
- What problems did you encounter while solving Exercise 12?

Your Expectations

https://www.klicker.uzh.ch/algodat

Student's Difficulty Level Assessment of Previous Assignments

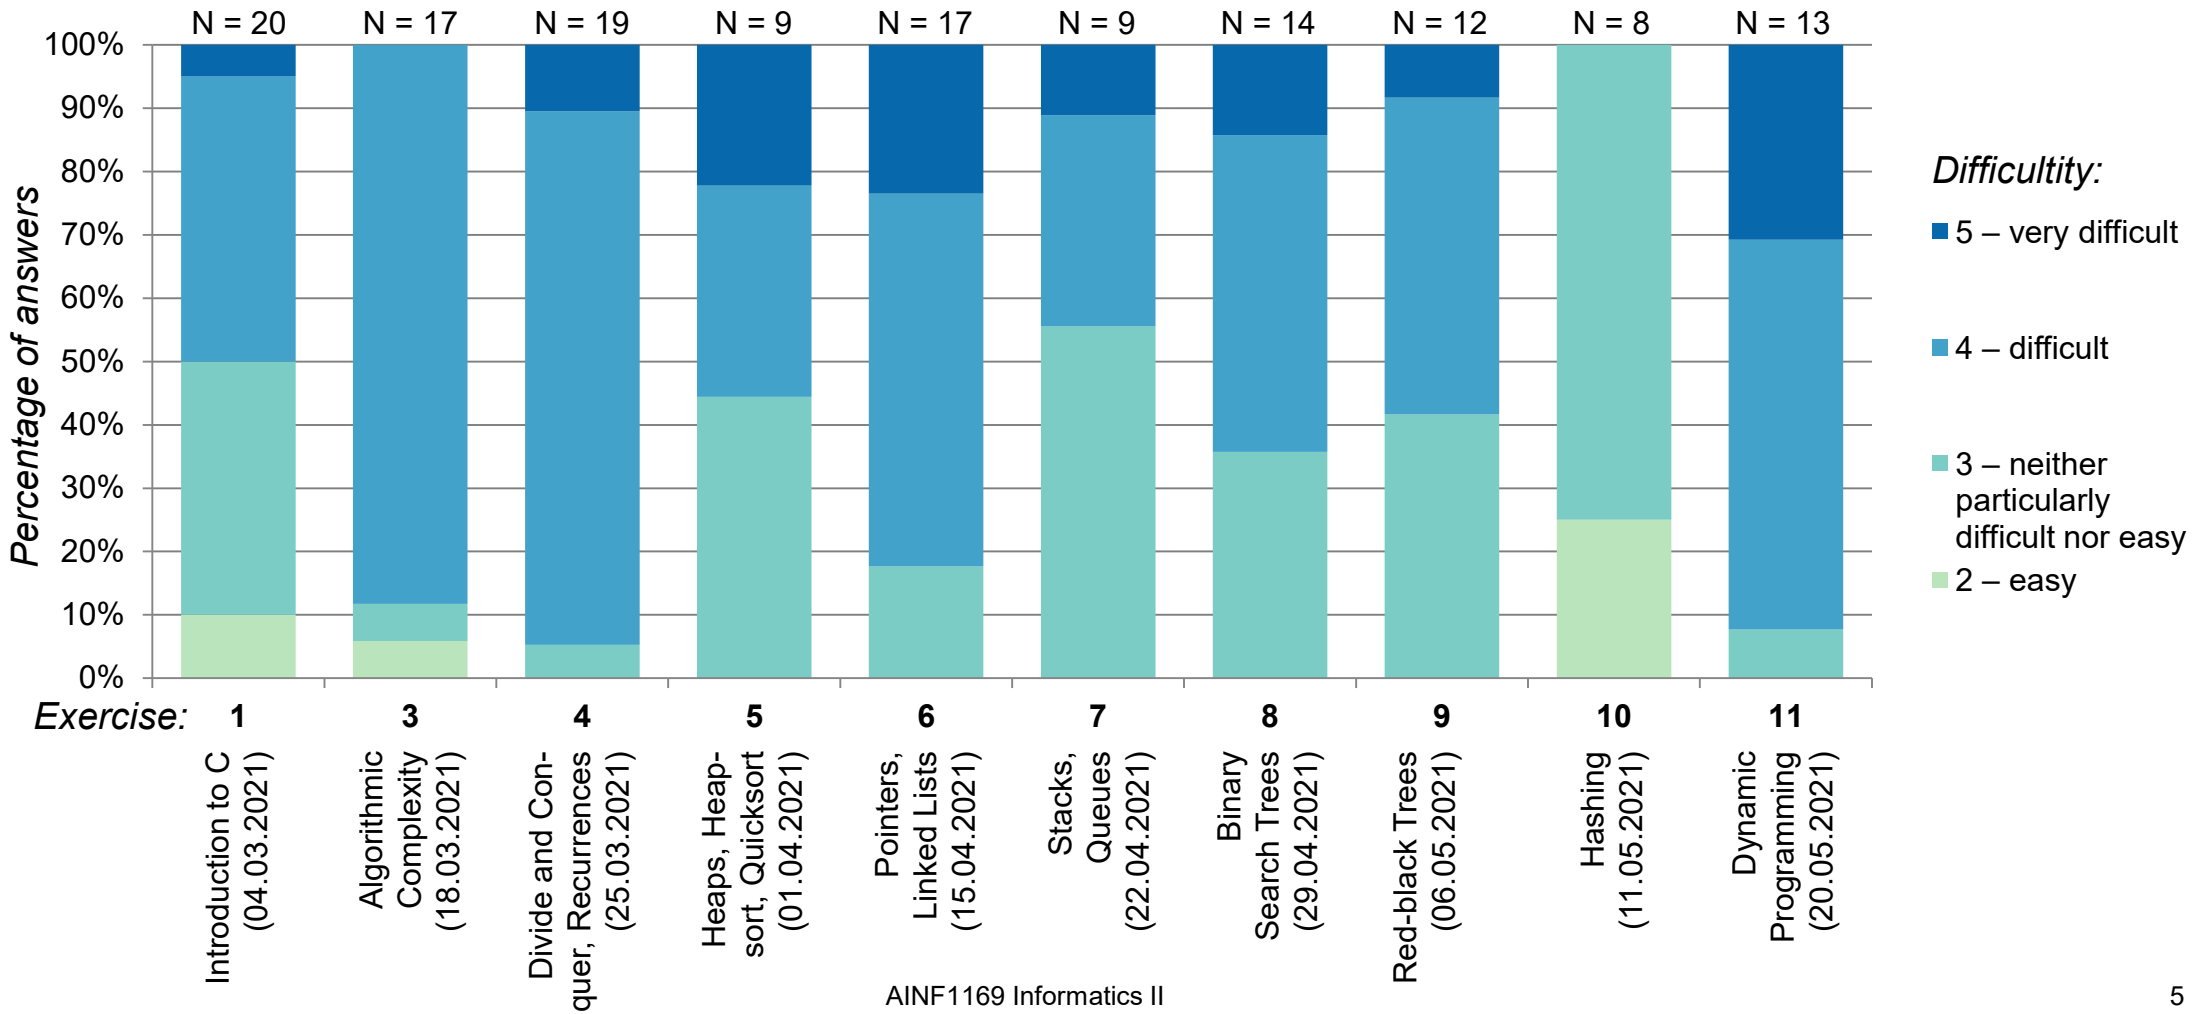

Midterm 1 Results

- Grading Scheme, Grading Curve
- Results Distribution

Midterm 1 Results

Not considering grades < 1.0:

- Average: 4.11
- Median: 4.25
  - Standard deviation: 0.96
  - Skewness: −0.84
  - Kurtosis: +0.05

- Registered on OLAT: 290
  - Grade 1.0: 49 (16.9 %)
  - Grade above 1.0: 241
- Passing grades: 167 (69.3 %)
- Failing grades: 74 (30.7 %)

- Reachable points: 40 pt
- Passing grade: ≥ 11.5 pt
  - Two best results: 27.5 pt, 27.0 pt (once each), grade 5.75

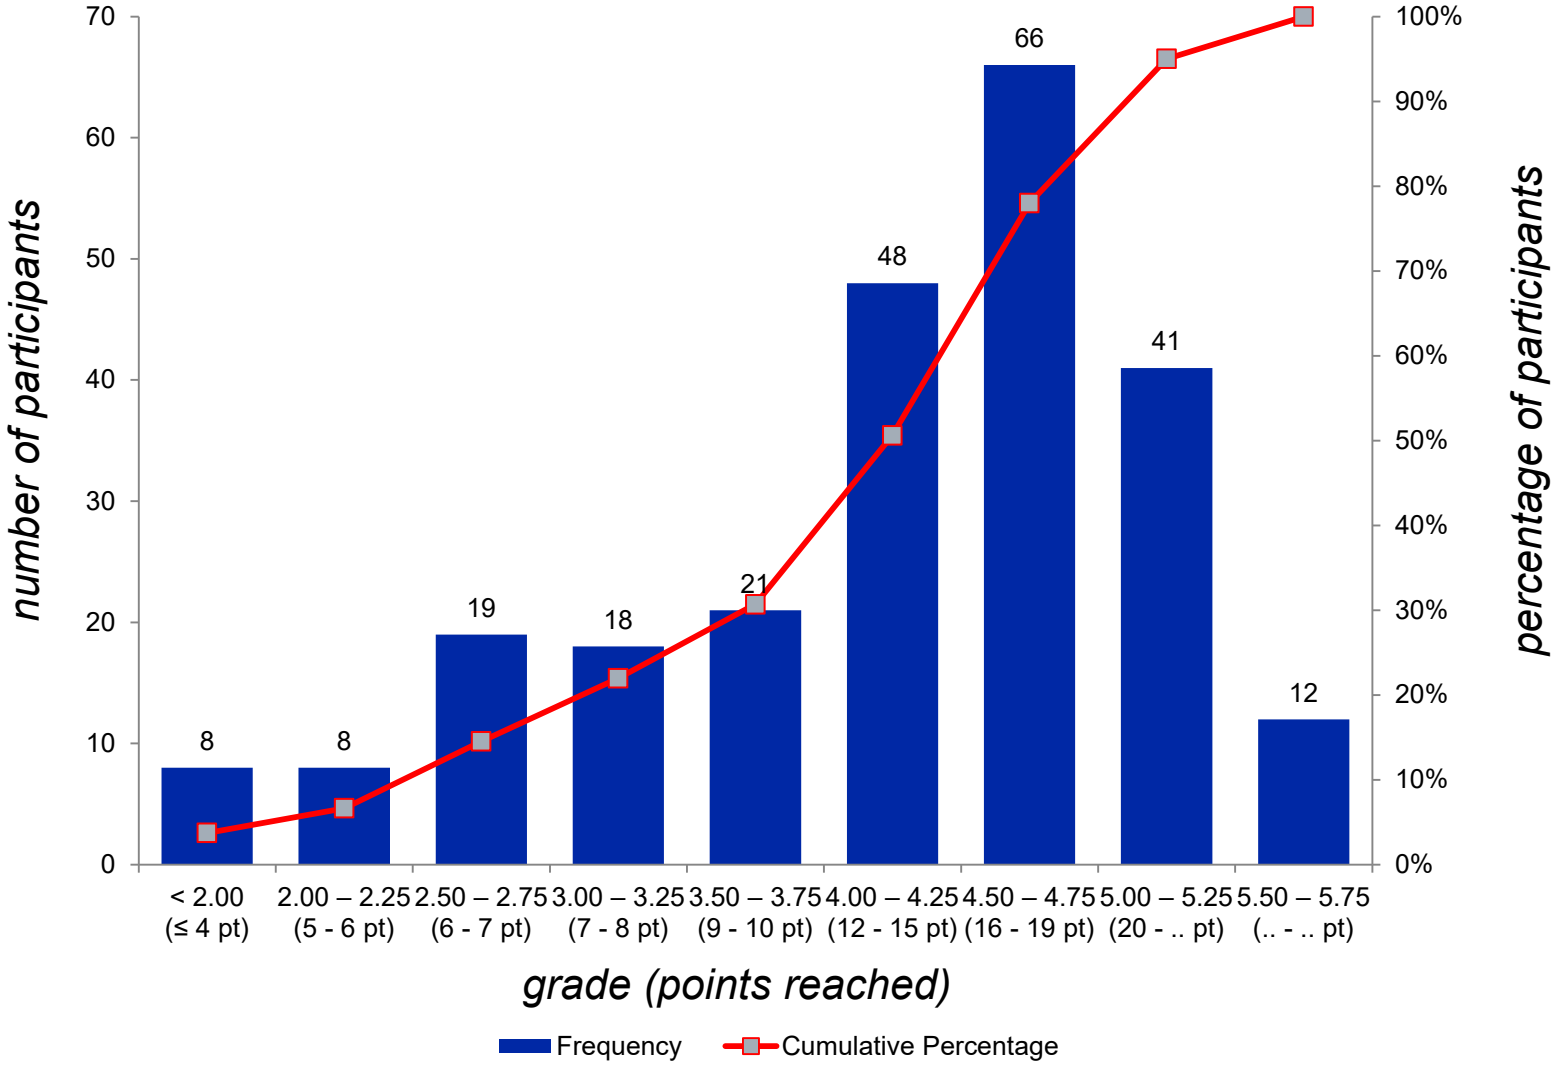

Administrivia

- Unofficial Q&A session

Unofficial Q&A Session

For people who want to ask questions, I will be available

tomorrow Friday, 28th of May

from 14.00 h

on Zoom

This will be an unofficial and informal Q&A session. I will not prepare any slides or materials but just (try to) answer any questions. The session will be recorded if people agree. If possible, send me any questions in advance.

Graph Algorithms

- Representation of Graphs: Adjacency Lists, Adjacency Matrices
- Graph Traversals: BFS and DFS
- Topological Sorting
- Minimum Spanning Trees, Prim-Jarník Algorithm
- SSSP, Dijkstra's Algorithm, Bellman-Ford Algorithm

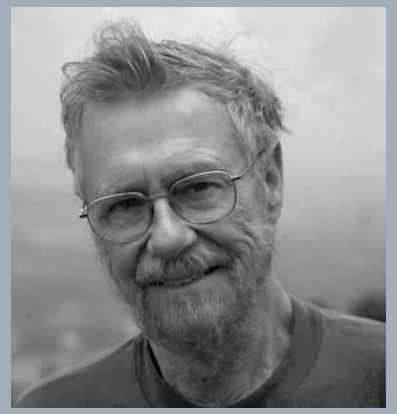
Edsger Dijkstra

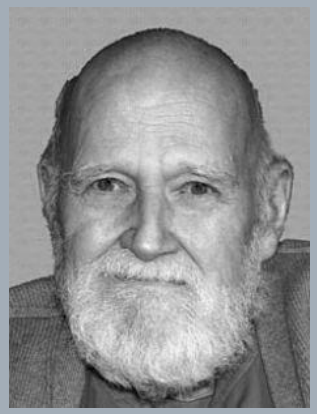
Lester Ford

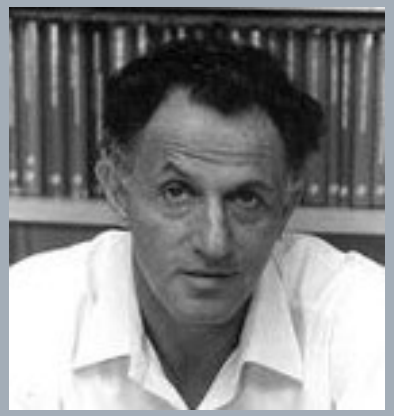
Richard Bellman

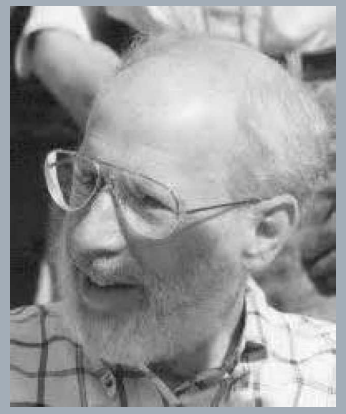
Joseph Kruskal

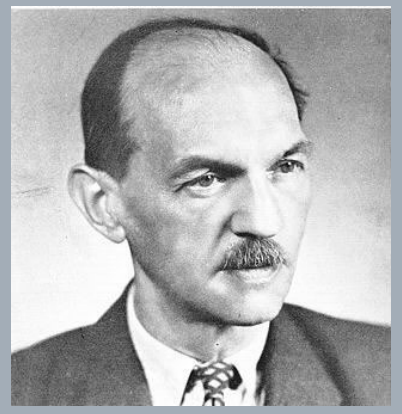
Vojtěch Jarník


Robert Prim

Graphs: Nomenclature and Different Types

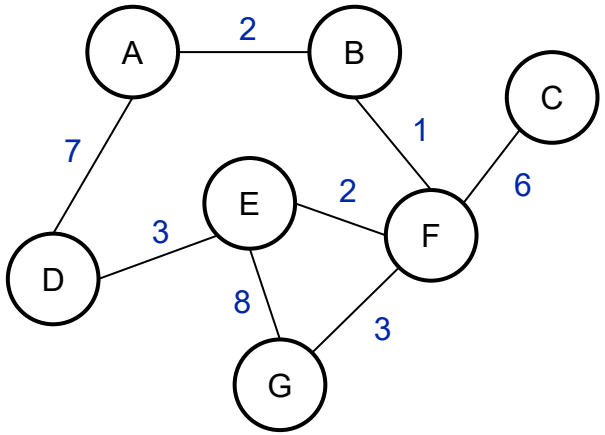

Some of the most important properties to describe graphs are:

- directed / undirected
- weighted / unweighted
  - with / without negative weights
- connected / disconnected
- with / without cycles (acyclic)
- with / without self-edges
- with / without multiple edges between to nodes (multigraph)

A given graph can fulfill one or more of these properties, producing a whole zoo of different kinds of graphs. For example a particularly important class are directed acyclic graphs (DAGs).

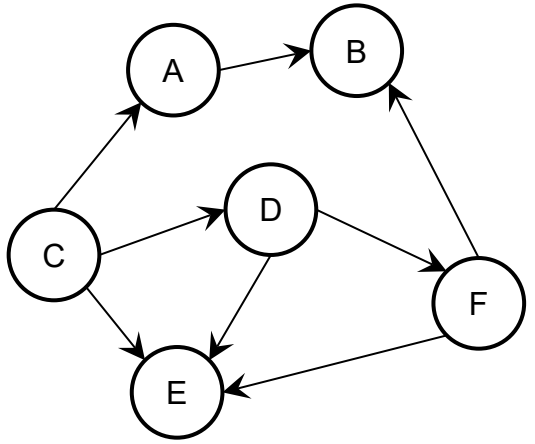

Representations of Graphs

Adjacency list

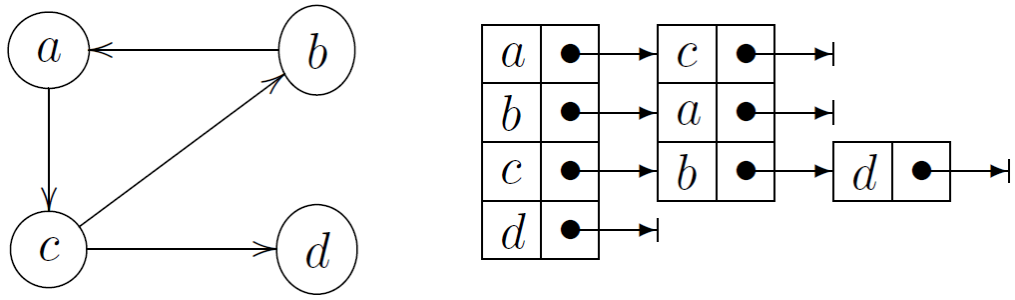

Adjacency matrix

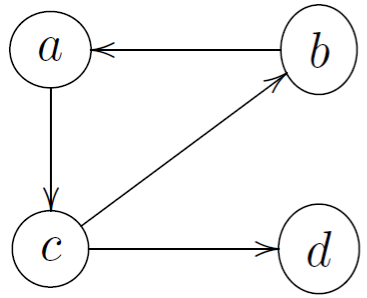

startpoint node \ *endpoint node*	a	b	c	d
a	0	0	1	0
b	1	0	0	0
c	0	1	0	1
d	0	0	0	0

- Exercise: Give a graph with five nodes which has an adjacency list where each list has a length of exactly one.
- Comprehension questions:
  - How to distinguish between directed and undirected graphs from looking at those representations? How could such a test be implemented?
  - How would you implement a function which finds the out-degree (number of outgoing edges) for a given node? What about the in-degree?

Exercise: Adjacency List vs. Adjacency Matrix

For each of the following cases, decide whether you would use the adjacency list structure or the adjacency matrix structure. Justify your choices.

- **A:** The graph has 10'000 vertices and 20'000 edges, and it is important to use as little space as possible.
  adjacency list: The adjacency list will use 10'000 + 20'000 = 30'000 units of space whereas the adjacency matrix will use $(10'000)^2$ = 100'000'000 units of space.
- **B:** The graph has 10'000 vertices and 20'000'000 edges, and it is important to use as little space as possible.
  adjacency list: The adjacency list will use 10'000 + 20'000'000 = 20'010'000 units of space whereas the adjacency matrix will use $(10'000)^2$ = 100'000'000 units of space.
- **C:** The graph has 10'000 vertices and you need to answer the query «isAdjacentTo» as fast as possible, no matter how much space you use.
  adjacency matrix: Each list of neighours in the adjacency list has a length of at most 10'000, so a query «isAdjacentTo» requires a scan through a list of this length in the worst case. In contrast, the respective query using a adjacency matrix takes constant time.

Adjacency Matrices vs. Adjacency Lists: Asymptotic Complexities

	Standard adjacency list (linked lists)	Fast adjacency list (hash tables)	Adjacency matrix
Space	$\Theta(V+E)$	$\Theta(V+E)$	$\Theta(V^2)$
Test if $uv \in E$	$O(1+\min\{\deg(u),\deg(v)\}) = O(V)$	$O(1)$	$O(1)$
Test if $u \to v \in E$	$O(1+\deg(u)) = O(V)$	$O(1)$	$O(1)$
List v's (out-)neighbors	$\Theta(1+\deg(v)) = O(V)$	$\Theta(1+\deg(v)) = O(V)$	$\Theta(V)$
List all edges	$\Theta(V+E)$	$\Theta(V+E)$	$\Theta(V^2)$
Insert edge uv	$O(1)$	$O(1)^*$	$O(1)$
Delete edge uv	$O(\deg(u)+\deg(v)) = O(V)$	$O(1)^*$	$O(1)$

Implementation of a Graph in C

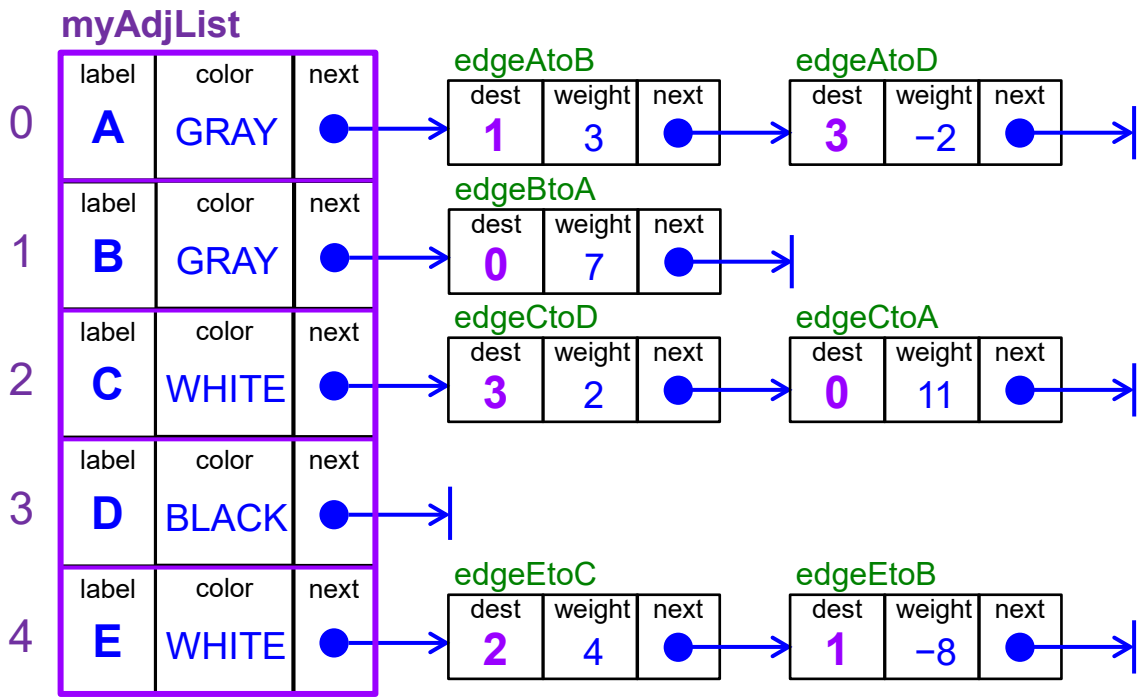


```
struct AdjacencyListNode {
    int destination;
    int weight;
    struct AdjacencyListNode* next;
};

struct GraphVertex {
    char label;
    int color;
    struct AdjacencyListNode* head;
};

struct Graph {
    struct GraphVertex* adjacency_list;
    int number_of_vertices;
};
```

Comprehension Question

Why do we represent graphs using adjacency lists or adjacency matrices? Why not just model them with structs which are interlinked with pointers as we do it with binary trees for example?

Lab Problem 13.1: Find Cycles, Check Connectivity

- Write a C program that checks whether a graph contains a loop.
- Write a C program that checks whether a graph is connected.

Graph Traversals: Introduction

There are two basic ways how the nodes in a graph can be visited:

- **Breadth-first search (BFS):** starting at some selected node, explore all neighboring nodes first and only if there are no unvisited neighbours left, proceed to do the same with the neighbours of the neighbours.
- **Depth-first search (DFS):** starting at some selected node, explore neighboring nodes as far as possible until a dead end (i.e. node with no more unvisited neighbours) is reached.

Both algorithms can be applied on directed or undirected graphs and will ignore any weights.

Remark: Tree traversals preorder, inorder and postorder are examples of depth-first search applied on a binary tree. A levelorder tree traversal is an example of a breadth-first search applied on a binary tree.

Breadth-First Search: Analogy

BFS expands like the waves around a stone thrown into a pond.

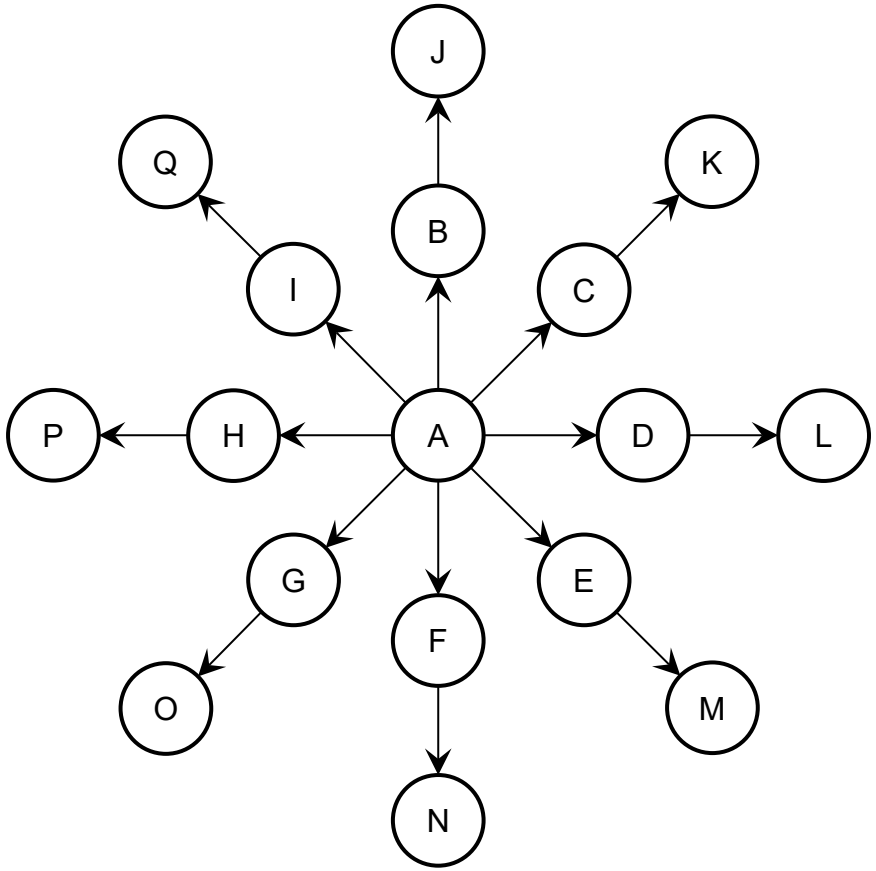

Depth-First Search: Principle

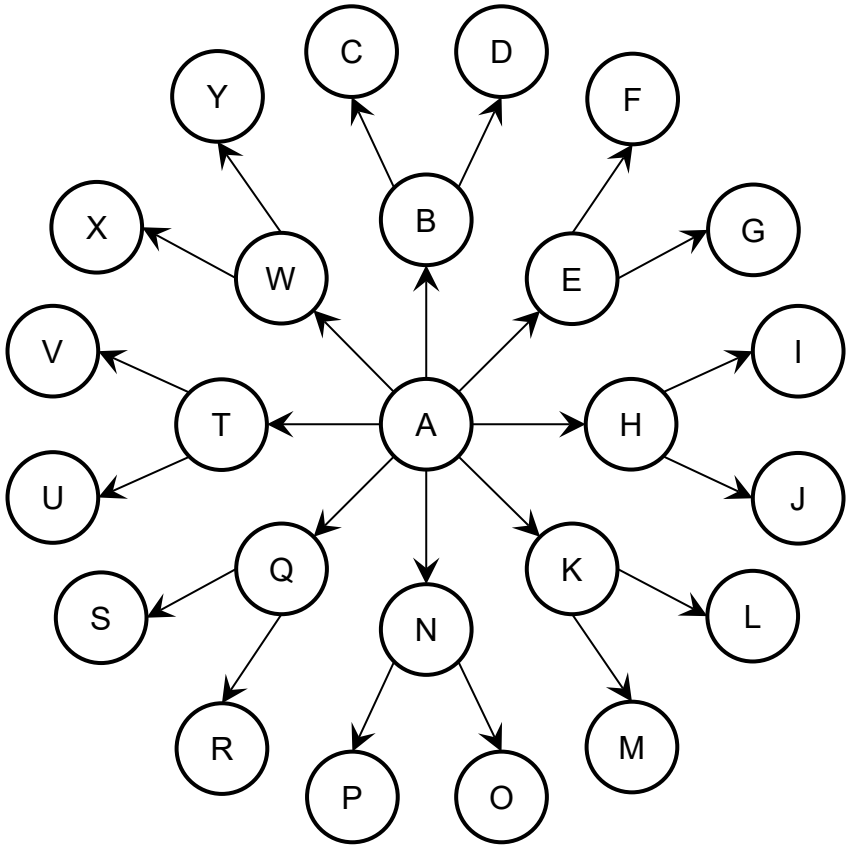

Graph Traversals: Book-Keeping Information

For some particular applications which we will see later, we will keep track of additional information while traversing the graph:

- For both BFS and DFS: color, predecessor
- BFS: distance (= number of edges traversed from starting node)
- DFS: start time, end time

Breadth-First Search: Example

– What is the order in which nodes are visited in the following graph, when applying BFS and starting at node A?

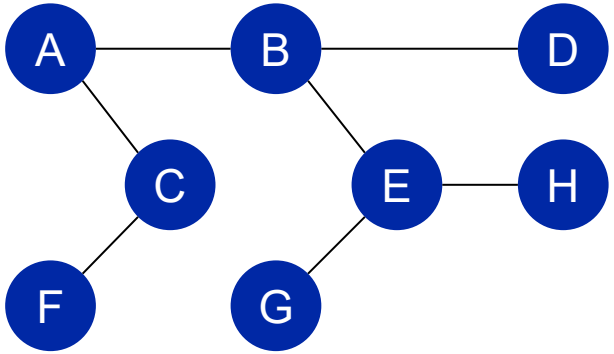

– Comprehension question: What other name could you assign to this particular way of passing through this graph?

BFS Algorithm

Algo: BFS(G,s)

```
foreach v ∈ G.V do
    v.col = W; v.dist = ∞; v.pred = NIL
s.col = G; s.dist = 0;
InitQueue(Q);
Enqueue(Q,s);
while Q ≠ ∅ do
    v = Dequeue(Q);
    foreach u ∈ v.adj do
        if u.col==W then
            u.col = G; u.dist = v.dist+1; u.pred = v; Enqueue(Q,u)
    v.col = B;
return (m,c);
```

Depth-First Search: Example

– What is the order in which nodes are visited in the following graph, when applying DFS and starting at node A?

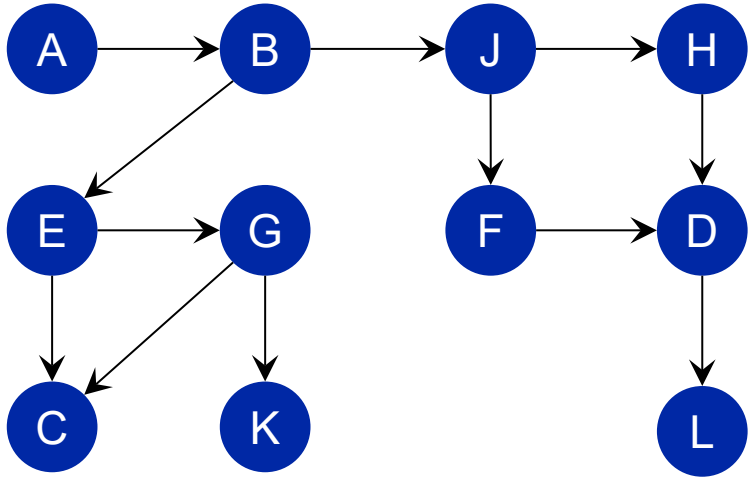

DFS Algorithm

Algo: DFS-Tree(G,s)

```
s.col = G; time = time+1; s.time = time;
foreach u ∈ s.adj do
    if u.col==W then
        u.pred = s;
        DFS-Tree(G,u)
s.col = B; time = time+1; s.etime = time;
```

DFS Algorithm

Algo: DFS(G)

```
foreach v ∈ G.V do
    v.col = W; v.pred = NIL
time = 0;
foreach v ∈ G.V do
    if v.col==W then DFS-Tree(G,v);
```

Breadth-First Search and Depth-First Search

	Breadth-first search	Depth-first search
Principle	Considers all directions	Persues one direction
Visiting	Might not visit all nodes of a graph; visits all node of a connected subgraph	Guarantees that all vertices are visited, regardless whether they are connected
Results	Single tree (distance of each visited node to the starting node)	Set of trees (start time, end time, edge classification)
Data structure	Queue	Stack
Complexity	O(|V| + |E|)	O(|V| + |E|)
Applications	Finding shortest paths	Cycle detection, topological sort, planarity test

Depth-First Search Variants

Algo: DFS-explicitStack(graph, startVertex)

```
1  S = initializeEmptyStack();
2  foreach u ∈ graph.vertices do
3      u.color = WHITE;
4  S.push(startVertex);
5  startVertex.color = GRAY;
6  while S.empty() == false do
7      u = S.peak();
8      if u.adjacentVertices == ∅ then
9          u.color = BLACK;
10         S.pop();
11     else
12         w = selectOneFrom(u.adjacentVertices);
13         w.color = GRAY;
14         S.push(w);
```

Algo: DFS-iterative(graph, startVertex)

```
1  S = initializeEmptyStack();
2  foreach u ∈ graph.vertices do
3      u.visited = false;
4  S.push(startVertex);
5  while S.empty() == false do
6      u = S.pop();
7      if u.visited == false then
8          u.visited = true;
9          foreach w ∈ u.adjacentVertices do
10             if w.visited == false then
11                 S.push(w);
```

Exercise 12 – Task 1: Breadth-First Search By Hand

In the given graph below, each vertex has an unique label. For example, we use vertex A to denote the vertex with label ”A”. Write a breadth first search (BFS) starts at vertex A using a queue. In this task, during the BFS search, neighbors of a vertex are visited in the alphabetical order of their labels. For example, in the BFS that starts at vertex A, vertex B is visited before vertex C. The first two steps of the solution are shown below.

…

Exercise 12 – Task 2: Depth-First Search By Hand

Given a graph below, each vertex has an unique label. Write the depth first search (DFS) starts at vertex A using a stack. In this task, during the DFS search, neighbors of a vertex are traversed in the alphabetical order of their labels. For example, in the DFS that starts from vertex A, vertex B is visited before vertex C. Note that the recursive solution of DFS is different from the one using a stack.

…

Spanning Trees

A spanning tree of a graph is a tree (a graph with no cycles) which contains all nodes of a graph.

Both BFS and DFS traversals of a graph produce (by considering the edges which were used for visiting the nodes) spanning trees (or potentially a set of trees, a forest, for unconnected graphs or for directed graphs in certain situations).

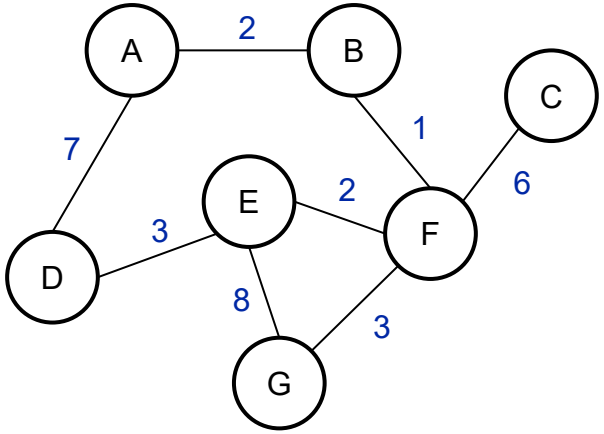

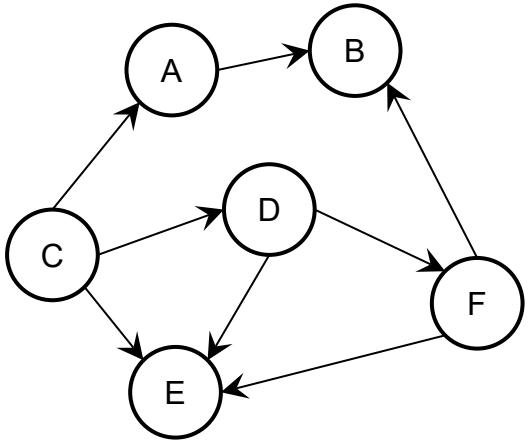

Spanning Trees

A spanning tree of a graph is a tree (a graph with no cycles) which contains all nodes of a graph.

Both BFS and DFS traversals of a graph produce (by considering the edges which were used for visiting the nodes) spanning trees (or potentially a set of trees, a forest, for unconnected graphs or for directed graphs in certain situations).

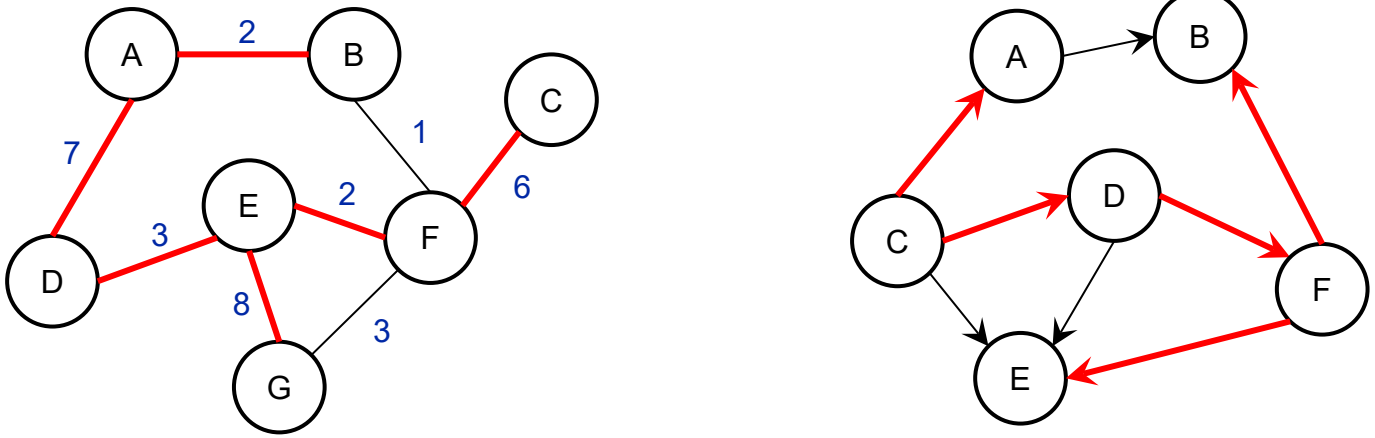

DFS: Additional Information

By applying the DFS algorithm, additional information about the nodes and edges in a graph can be gathered and stored:

- start time and end time of each node
- edge classification for each edge

This information will be valuable for various purposes as we will see later on.

DFS: Start Time and End Time

While performing the DFS traversal of a graph, we can note down for each node the occurence of two events:

- Start time: when we first visited a node (encountered / started a node, made it grey, put the node in the stack)
- End time: when we last visited a node (finished a node, made it black, popped it from the stack)

Remark:

- «time» is a bit odd since not any actual measured time in the physical sense is meant but just the sequence / position / counter

Edge Classification with DFS

With regard to the spanning tree produced by the DFS algorithm, we can distinguish four different types of edges:

- tree edges: from a parent to a child, used for visiting nodes in DFS algorithm, edges forming the spanning tree produced by DFS algorithm
- forward edges: to a non-child descendant / to a finished vertex discovered after the current vertex
- back edges: to an ancestor / to a discovered but unfinished node
- cross edges: everything else (i.e. to a vertex finished before the current vertex's discovery); cross edges connect independent subtrees

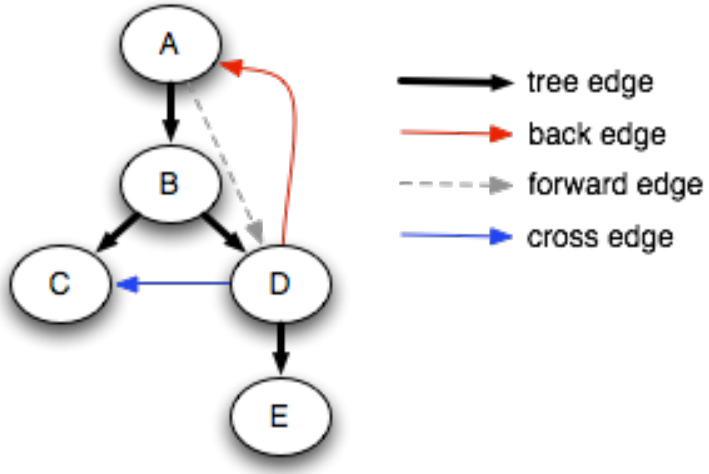

Examples / Exercises Edge Classification

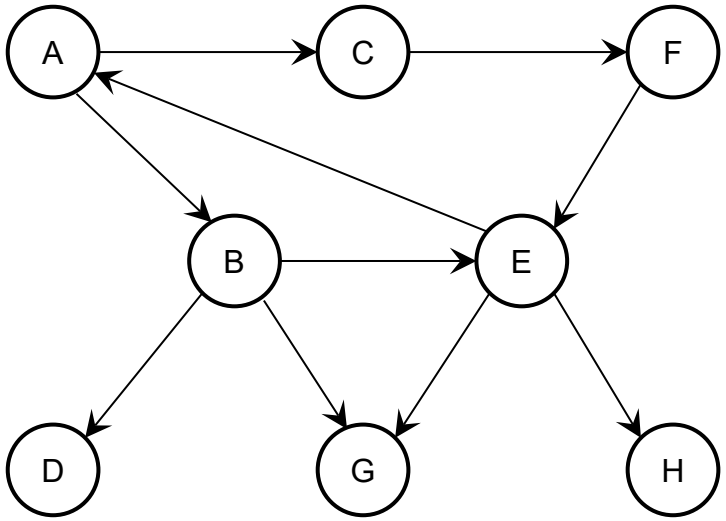

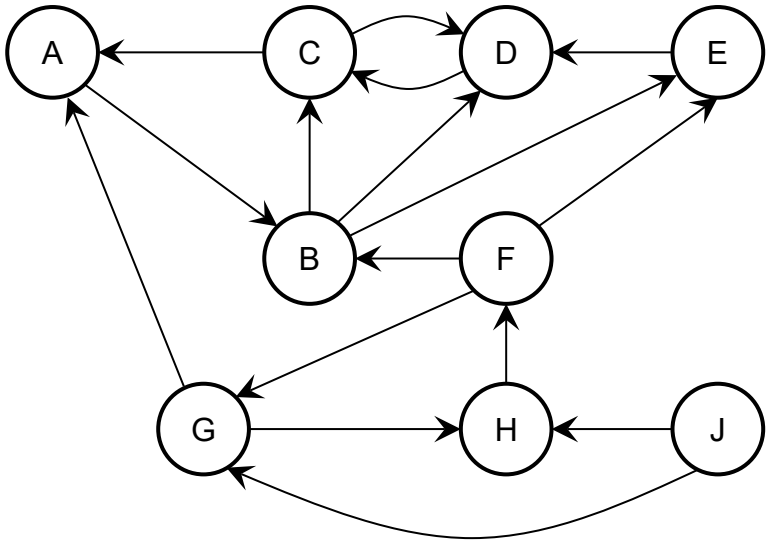

Examples / Exercises Edge Classification: Solution

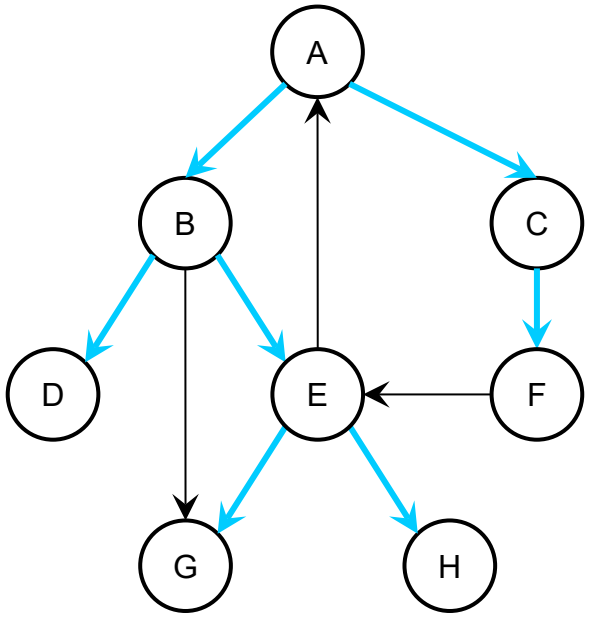

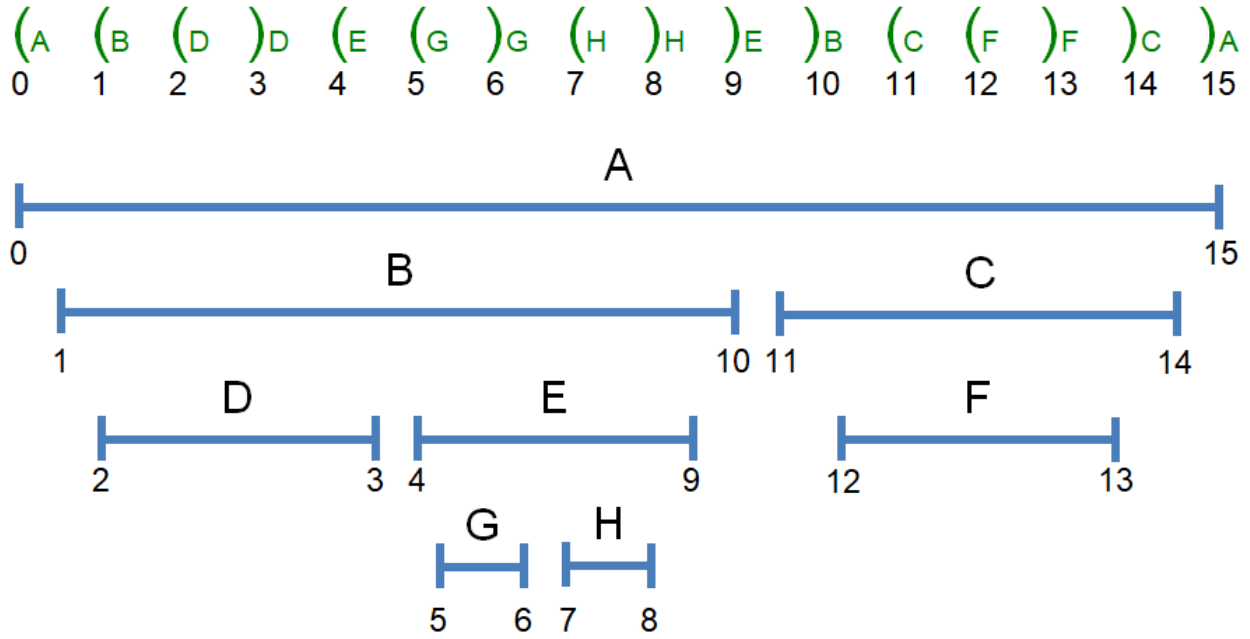

Edge Classification, Start and End Times

– How can edge detection be implemented? And how do edge classification and start / end time relate to eachother?

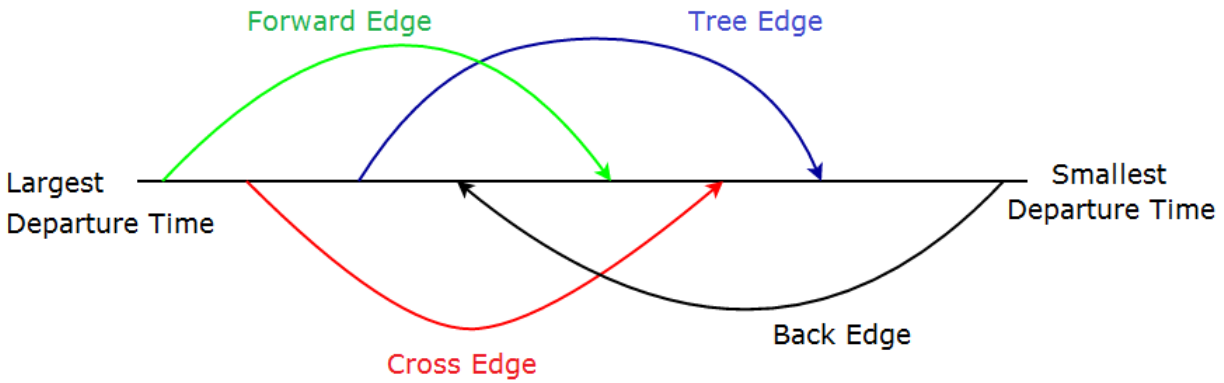

– These considerations also point to the parentheses theorem (and they are also important for topological sorting)…

Parentheses Theorem

This means that, for two nodes A and B, the intervals as given by the start and end times of a node, $[s_A, e_A]$ and $[s_B, e_B]$, must be in one of the following two configurations:

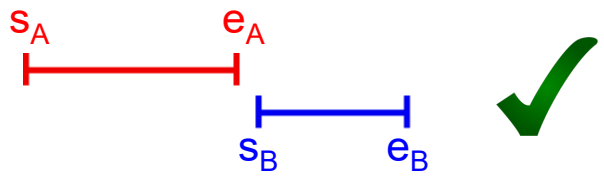

disjoint intervals

- The intervals of two nodes A and B are entirely disjoint.
- Start and end times of the second interval of node B are both encountered *after* the intervall of the first node A has finished.

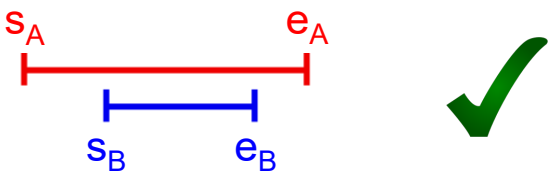

nested intervalls

- The intervals of one node (here: A) is contained entirely within the interval of another node (here: B).
- The inner interval B is encountered after A has started and B ends before A ends.

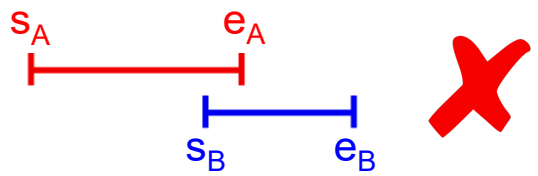

overlapping intervals

- The two intervals of two nodes overlap each other.
- The second interval B starts before the first A has ended but B finishes after A.
- This configuration cannot occur.

Parentheses Theorem

The parentheses theorem can be derived directly from the properties of the spanning tree.

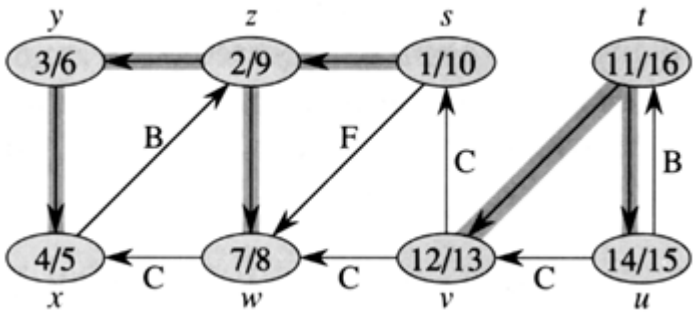

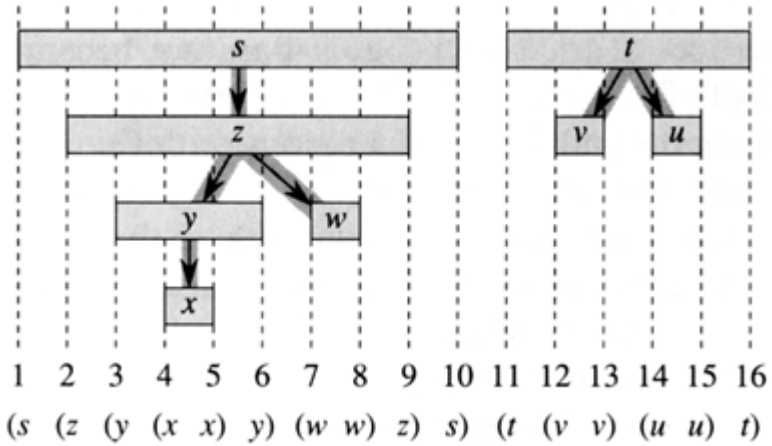

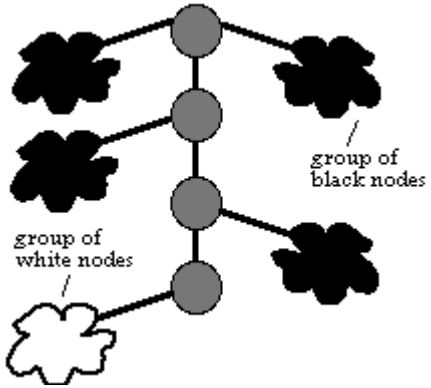

DFS Application: Cycle Detection

DFS provides an easy way to determine whether a given graph has a cycle: If the spanning tree produced by the DFS traversal contains a back edge, the graph contains a cycle.

Topological Sorting: Introduction and Example

- Any directed acyclic graph (DAG) has a topological order (and normally even many different ones).
- If a graph has a cycle, a topological order cannot exist.

Example: How to get dressed in the morning

- The adjacent graph shows the dependencies of which pieces of clothing have to be put on necessarily before other ones. For example, the socks have to be put on before the shoes are put on.
- In which order can the clothes be put on in accordance with this graph?

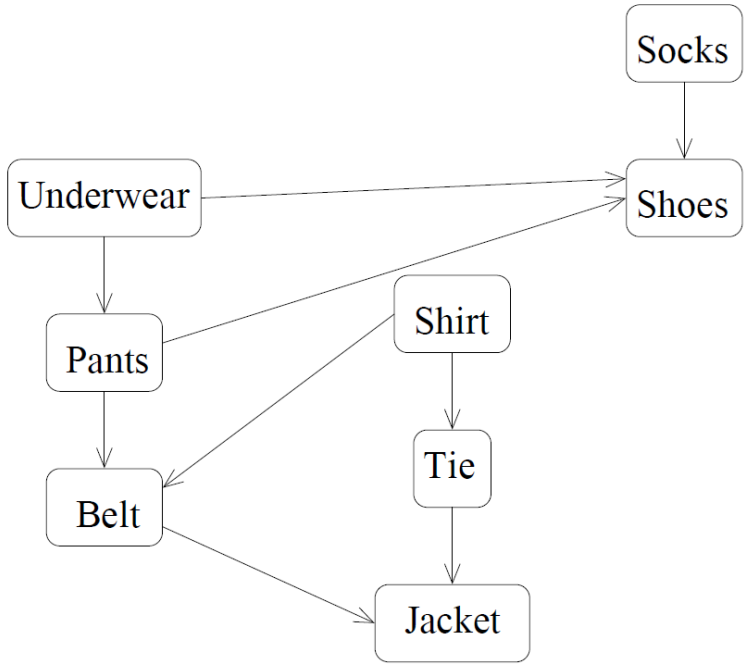

Topological Sorting: Example: How to Get Dressed in the Morning

– Possible solution:

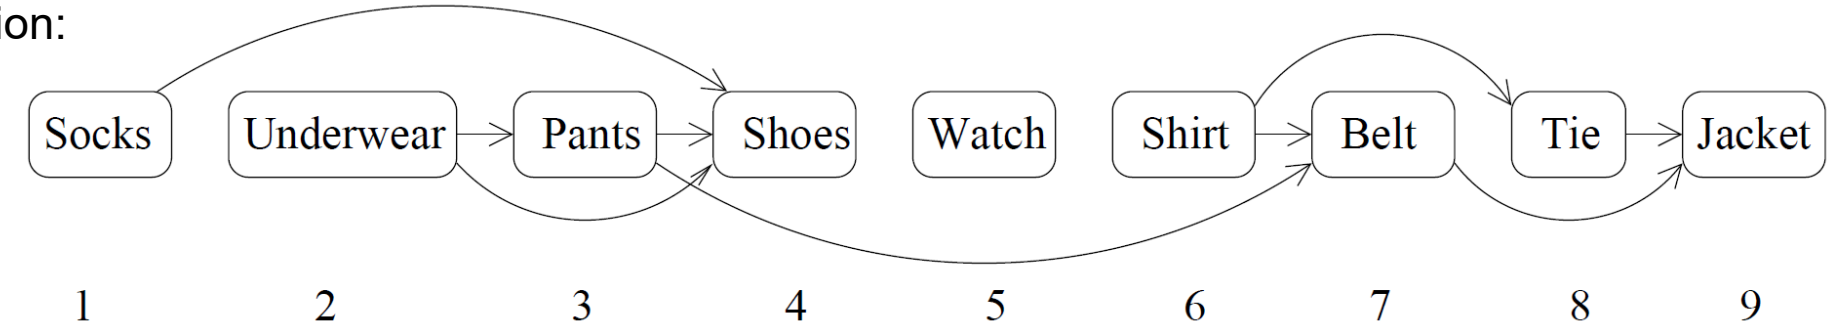

– Other solutions are possible.

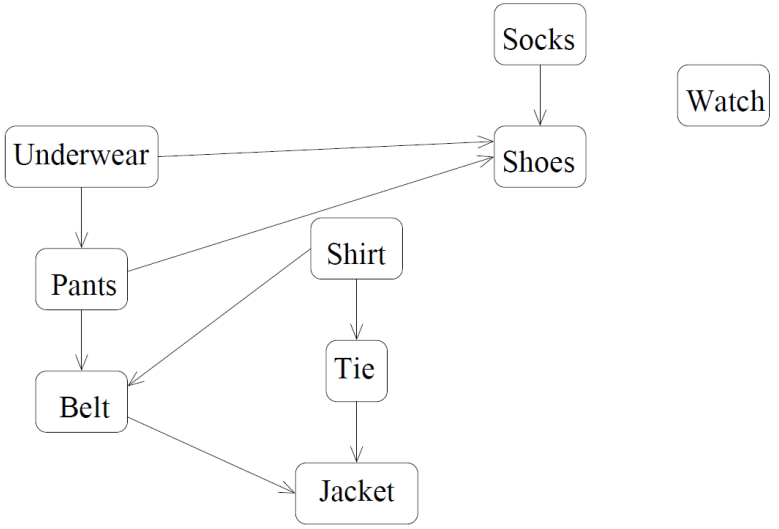

Topological Sorting: Implementation with DFS

Topological order of a graph G can be established using the depth-first search (DFS) algorithm:

– Use DFS(G) to get start and end times for all vertices in G.
– Return list of vertices in reverse order of end times.

Additional Exercise: Islands / Ecosystems

(Task 3 of assignment 6 from FS 2018)

You are given a binary 2D matrix M, each element of the matrix represents an island, islands with value 0 are uninhabited and islands with value 1 are inhabited. Each island represented by a cell in the matrix can have up to 8 neighbors (north, east, west, south and 4x diagonally). A group of neighboring inhabited islands form an ecosystem. The goal is to find the number of distinct ecosystems in the archipelago of islands represented by the 2D matrix M.

1	1	0	0	1
0	1	0	0	0
1	0	0	1	1
0	0	0	0	0
1	0	1	0	1

Additional Exercise: Islands / Ecosystems

(Task 3 of assignment 6 from FS 2018)

You are given a binary 2D matrix M, each element of the matrix represents an island, islands with value 0 are uninhabited and islands with value 1 are inhabited. Each island represented by a cell in the matrix can have up to 8 neighbors (north, east, west, south and 4x diagonally). A group of neighboring inhabited islands form an ecosystem. The goal is to find the number of distinct ecosystems in the archipelago of islands represented by the 2D matrix M.

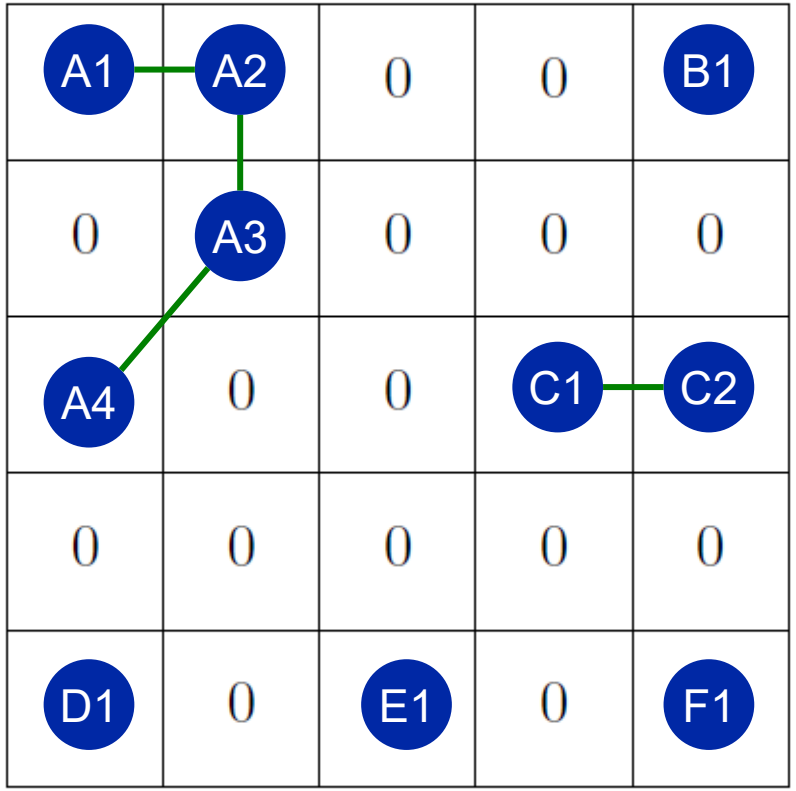

Idea for solution:

Call DFS on each grid position if there is a value of 1 and the node was not yet visited by another call of DFS. Count the number of times, DFS was called.

Additional Exercise: Islands / Ecosystems

```c
#include <stdio.h>
#include <string.h>

#define ROW 5
#define COL 5

static int M[ROW][COL] = {
    {1, 1, 0, 0, 1},
    {0, 1, 0, 0, 0},
    {1, 0, 0, 1, 1},
    {0, 0, 0, 0, 0},
    {1, 0, 1, 0, 1}
  };

int isSafe(int row, int col, int visited[ROW][COL]) {
  return (row >= 0) && (row < ROW) &&
         (col >= 0) && (col < COL) &&
         (M[row][col] && !visited[row][col]);
}
```

```c
void DFS(int row, int col, int visited[ROW][COL]) {
  int k;
  int rowNbr[] = {-1, -1, -1,  0, 0,  1, 1, 1};
  int colNbr[] = {-1,  0,  1, -1, 1, -1, 0, 1};
  visited[row][col] = 1;

  for (k = 0; k < 8; ++k)
    if (isSafe(row + rowNbr[k], col + colNbr[k], visited) ) {
      DFS(row + rowNbr[k], col + colNbr[k], visited);
    }
}
```

```c
int countEcosystems() {
  int i;
  int j;
  int visited[ROW][COL];
  memset(visited, 0, sizeof(visited));

  int count = 0;
  for (i = 0; i < ROW; ++i)
    for (j = 0; j < COL; ++j)
      if (M[i][j] && !visited[i][j]) {
        DFS(i, j, visited);
        ++count;
      }
  return count;
}
```

```c
void main() {
  printf("Number of ecosystems is: %d\n", countEcosystems());
}
```

Overview of Some DFS and BFS Applications

DFS applications:

- Cycle detection
- Topological sorting
- Connected components identification/counting
- (Planarity test)

BFS applications:

- Calculate the distance from a start node to each other (reachable) node (e.g.: Erdős number)
- (With extensions: finding shortest paths as part of Dijkstra's algorithm)

Exercise 12 – Task 3: DFS Implementation

This task is about implementing the depth-first search (DFS) on graphs. Graphs are represented by adjacency lists.

The following code snippet is everything you need for your graph:

…

Implement the function void DFS(struct node** graph, int start) that prints a DFS on the Graph from start vertex. Hint: You can find Advanced Data Structure implementations on previous exercise sheets.

Exercise 12 – Task 4: Maze Implementation

So far we have studied search in graphs. This task is about search in a 2D array (2D matrix). Imagine a 2D array, where each element shows a physical location in a maze. In this exercise, we use the following notation:

…

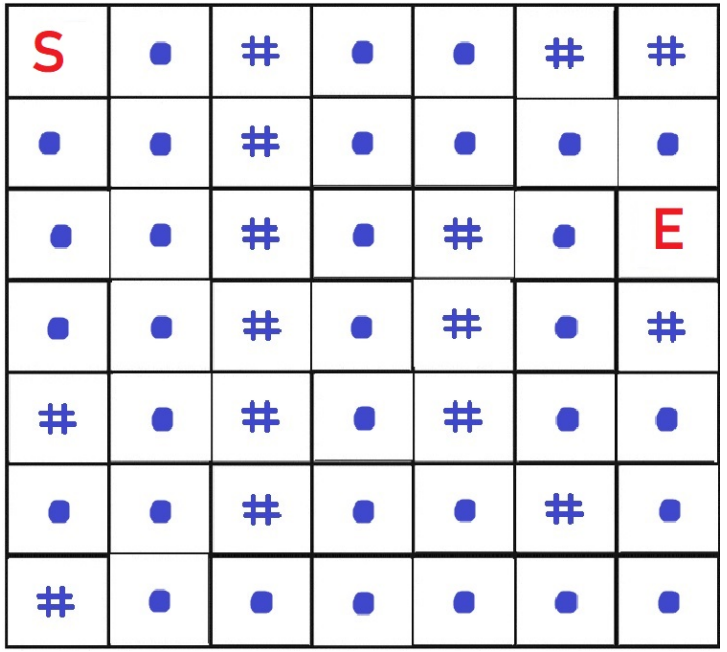

Break – Break: Warning

Note: All the topics contained in the slides from here on are part of chapter SL10 in the lecture slides and thus will not be part of the final exam.

Algorithm of Prim and Jarník

– Finds a minimum spanning tree (MST) for a given graph. Operates on weighted trees (directed or undirected).

– *Steps:*

(1) Select an arbitrary node as starting node.

(2) Repeatedly follow the least costly edge from the set of nodes which have already been visited until all nodes have been visited.

– The algorithm grows a tree from the starting node. It is a node oriented algorithm.

– Remark: There is another algorithm for this purpose called Kruskal's algorithm which instead creates the minimum spanning tree by selecting edges with the smallest weight from the set of all edges (not considering whether their nodes have already been visited). Kruskal's algorithm is edge oriented.

Prim-Jarník Algorithm: Example / Exercise

Considering the adjacent, simplified distance map of Switzerland, use the algorithm of Prim and Jarník to find a minimum spanning tree (MST).

Additional questions:

- What could be a possible use case for a minimum spanning tree like the one created in this task?
- Will the resulting MST always be the same, indepent on which node is used as start node (not only for this graph, but for arbitrary graphs)?

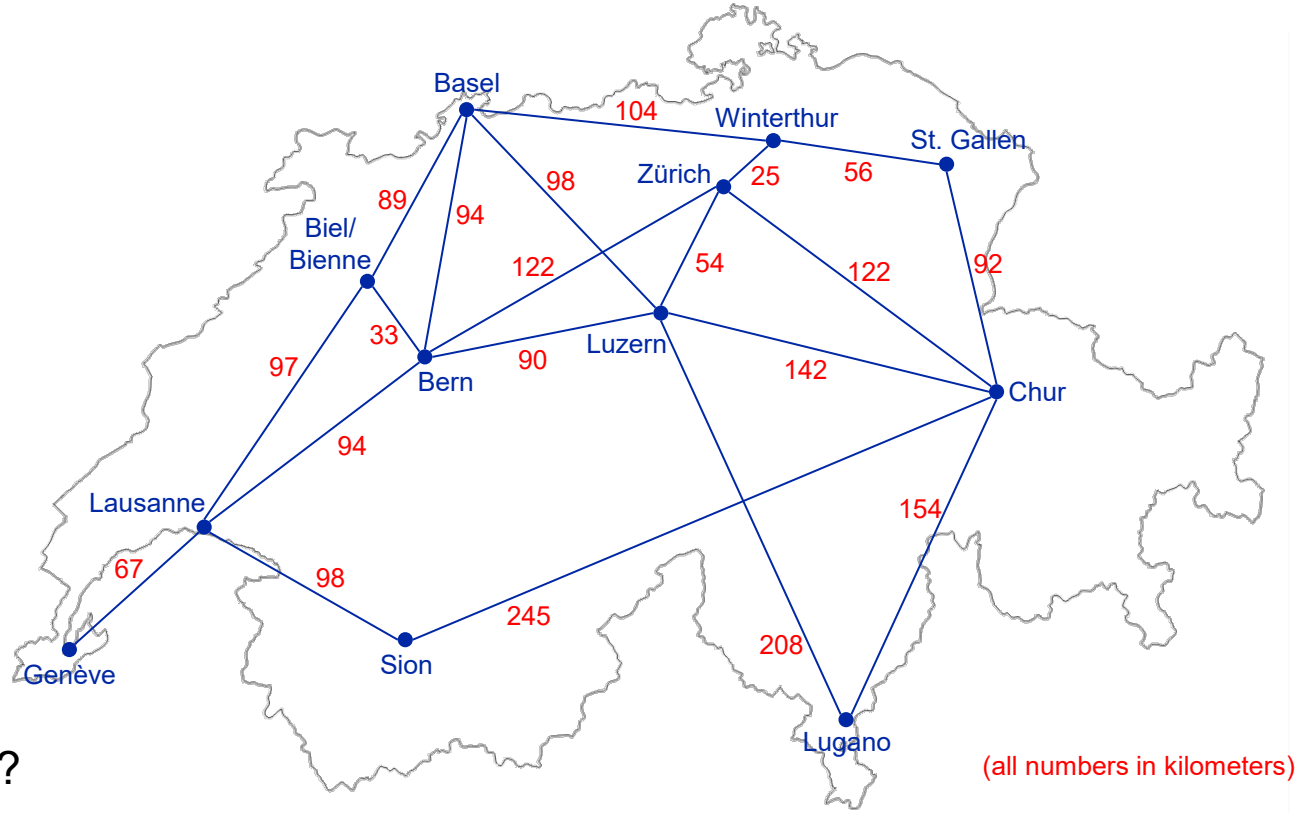

Prim-Jarník Algorithm: Example / Exercise

Considering the adjacent, simplified distance map of Switzerland, use the algorithm of Prim and Jarník to find a minimum spanning tree (MST).

Sequence of choosing the edges when starting in Zürich:

1) Zürich – Winterthur (25 km)
2) Zürich – Luzern (54 km)
3) Winterthur – St. Gallen (56 km)
4) Luzern – Bern (90 km)
5) Bern – Biel/Bienne (33 km)
6) Biel/Bienne – Basel (89 km)
7) St. Gallen – Chur (92 km)
8) Bern – Lausanne (94 km)
9) Lausanne – Genève (67 km)
10) Chur – Lugano (154 km)

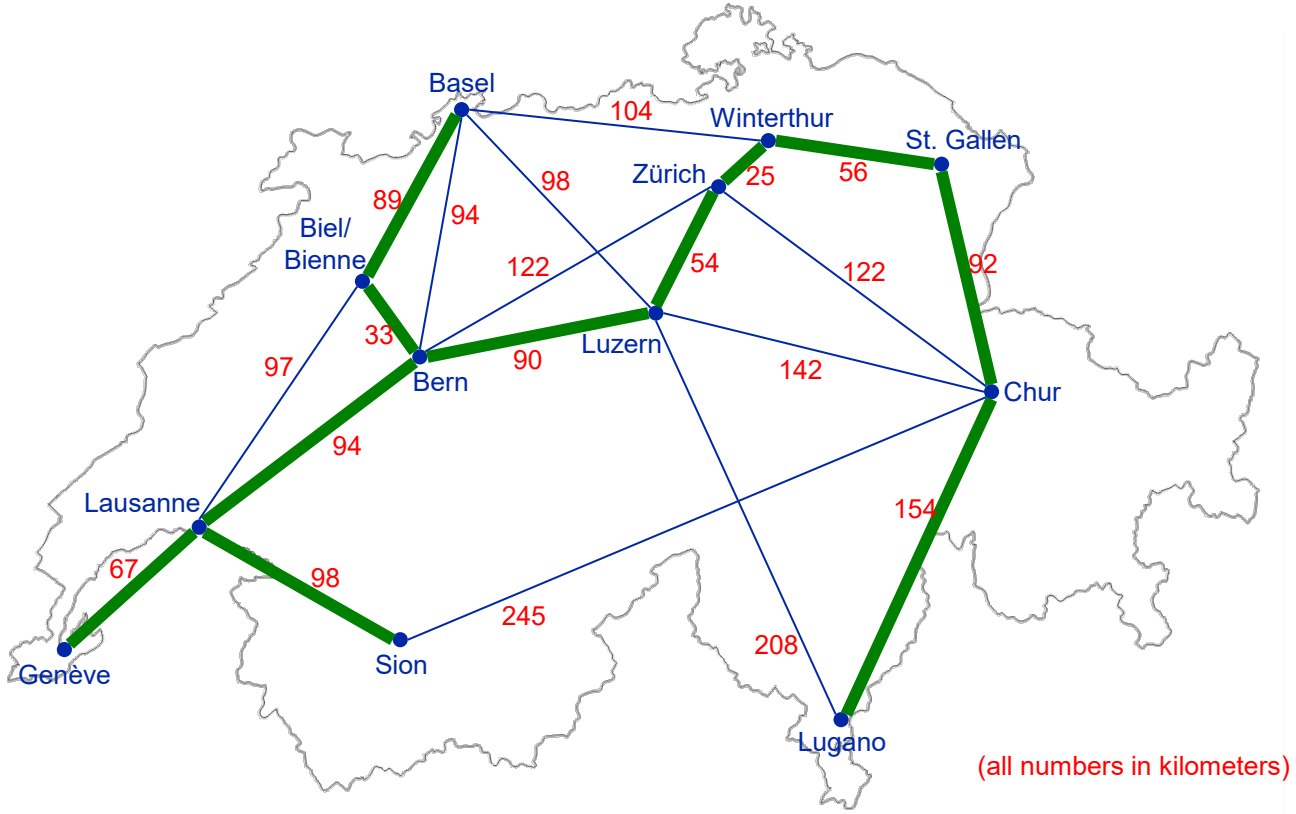

Prim-Jarník Algorithm Visualized

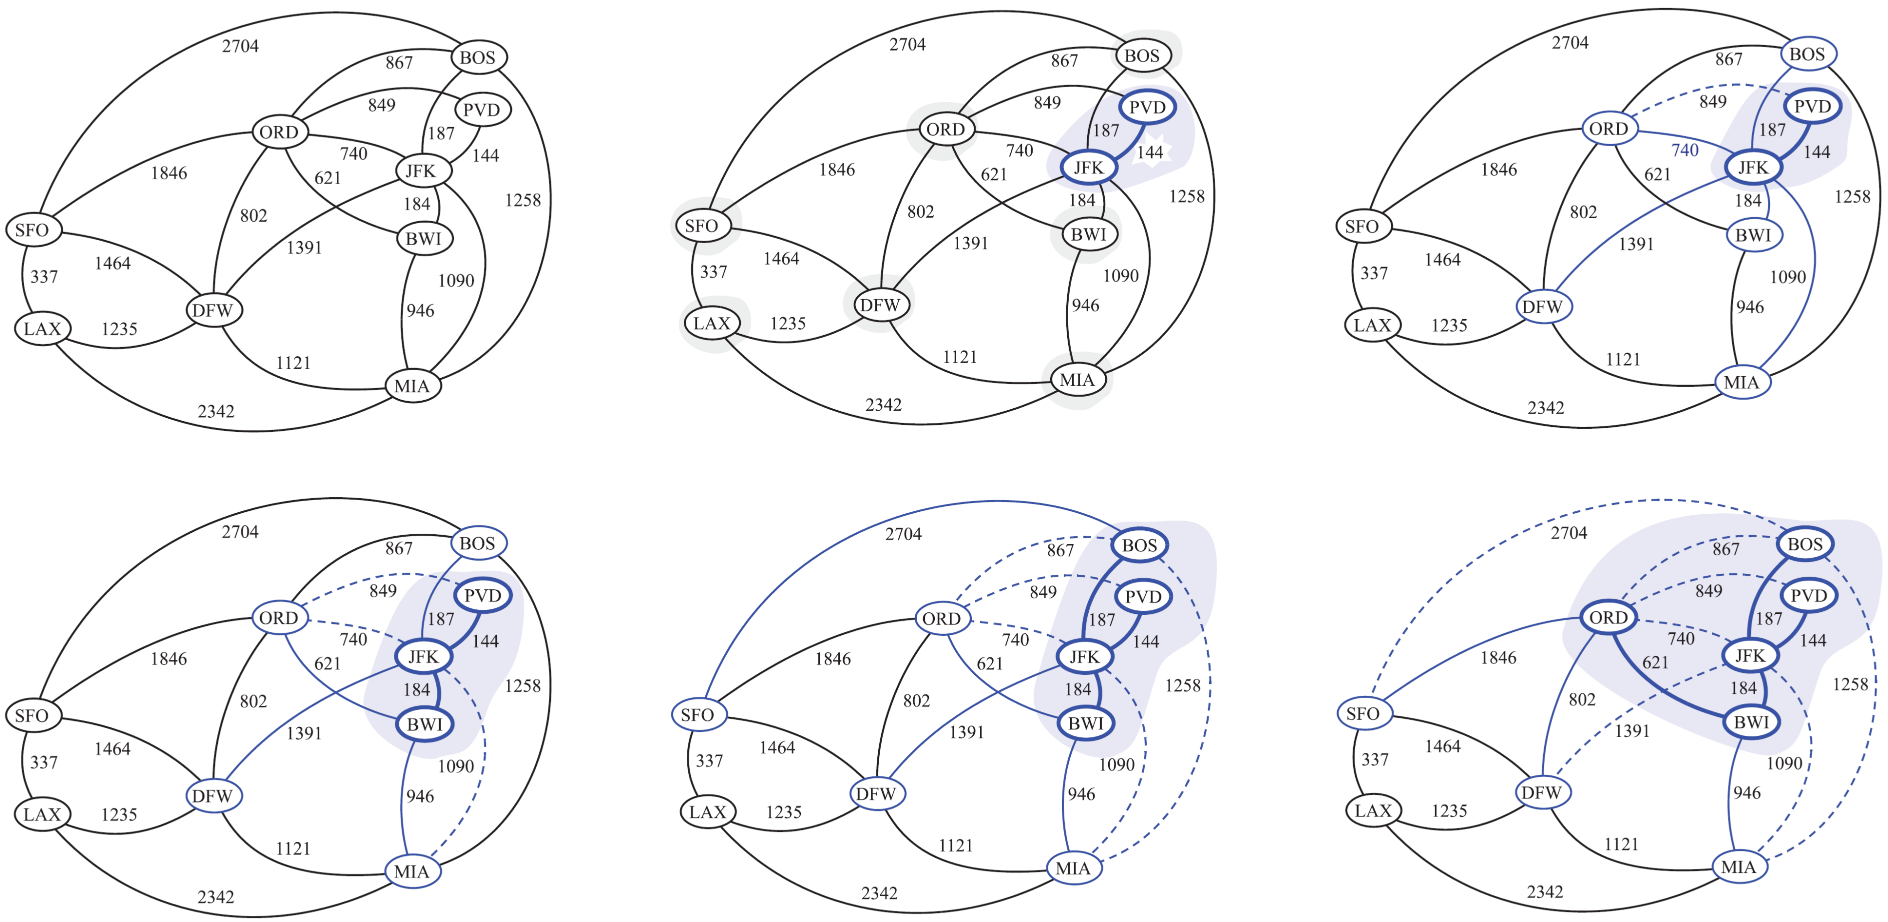

Prim-Jarník Algorithm Visualized

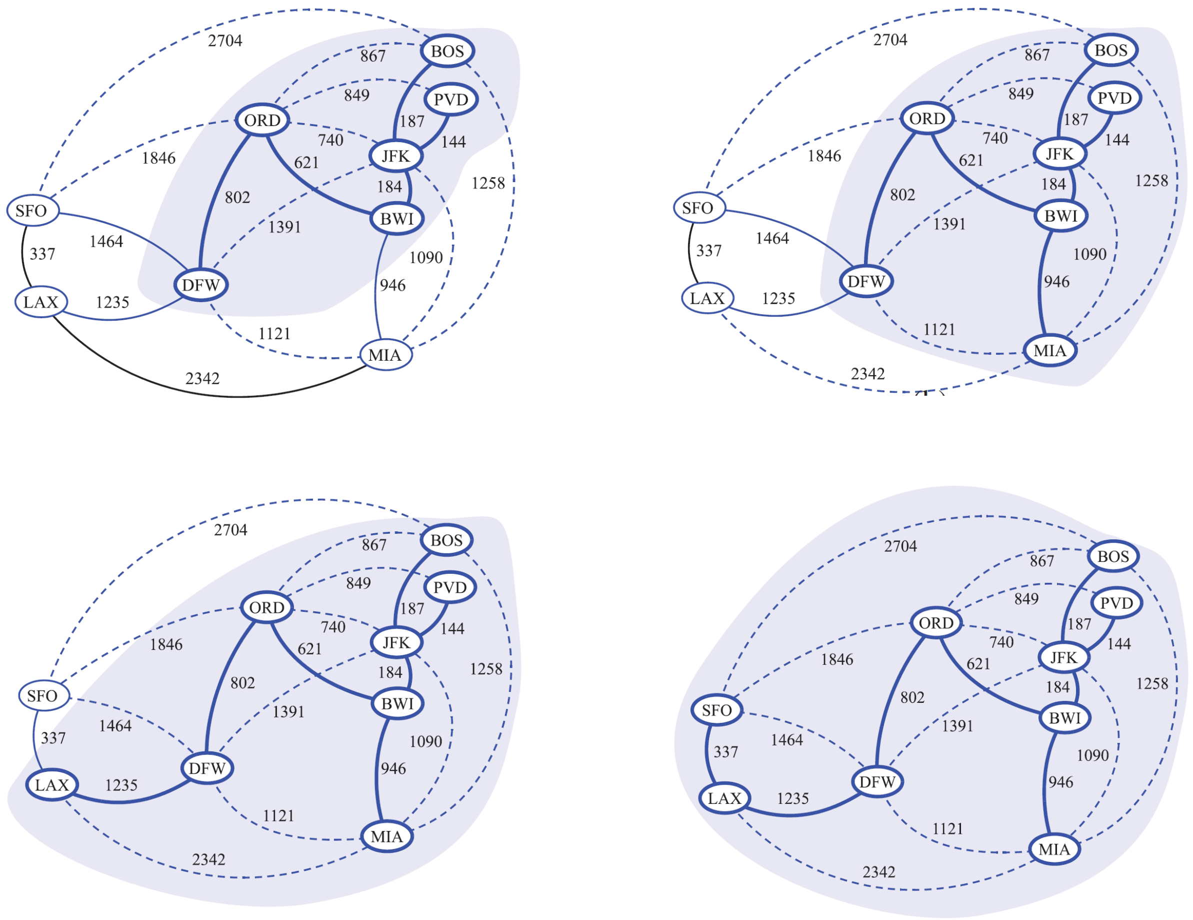

from: Michael T. Goodrich, Roberto Tamassia, David M. Mount: Data Structures and Algorithms in C++ (second edition, 2011). John Wiley & Sons; p. 652f.

Algorithm of Prim and Jarník: Implementation

Algo: PrimJarnik(G,w,s)

```
foreach v ∈ G.V do
    v.key = ∞;
    v.pred = NIL;
s.key = 0;
InitMinPriorityQueue(PQ,G.V);
while PQ ≠ ∅ do
    v = ExtractMin(PQ);
    foreach u ∈ v.adj do
        if u ∈ PQ ∧ w(v,u) < u.key then
            u.key = w(v,u);
            u.pred = v;
            DecreaseKey(PQ,u,w(u,v))
```

Kruskal's Algorithm: Example / Exercise

Considering the adjacent, simplified distance map of Switzerland, use Kruskal's algorithm to find a minimum spanning tree (MST).

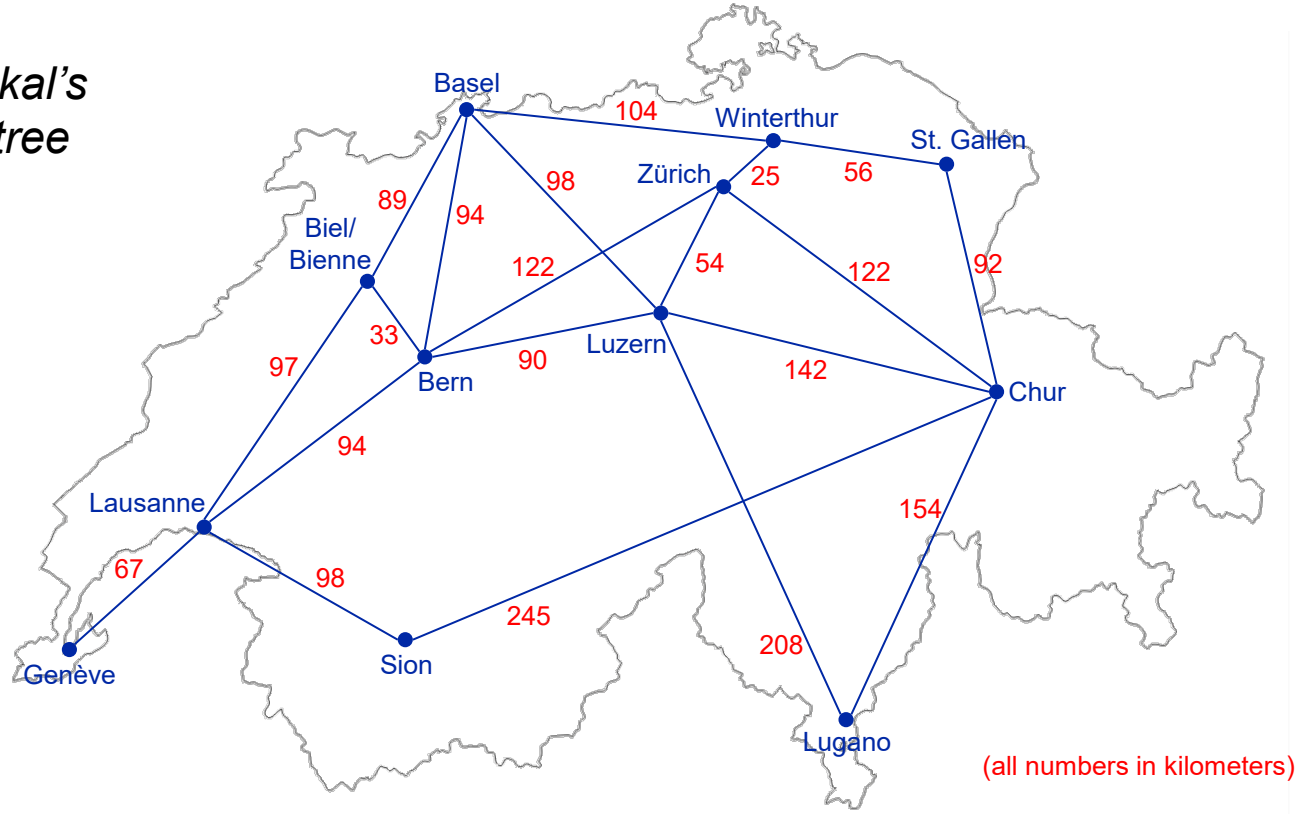

Kruskal's Algorithm: Example / Exercise

Considering the adjacent, simplified distance map of Switzerland, use Kruskal's algorithm to find a minimum spanning tree (MST).

Sequence of choosing the edges:

1) Winterthur – Zürich (25 km)
2) Bern – Biel/Bienne (33 km)
3) Winterthur – St. Gallen (56 km)
4) Genève – Lausanne (67 km)
5) Basel – Biel/Bienne (89 km)
6) Bern – Luzern (90 km)
7) St. Gallen – Chur (92 km)
8) Bern – Lausanne (94 km)
9) Lausanne – Sion (98 km)
10) Chur – Lugano (154 km)

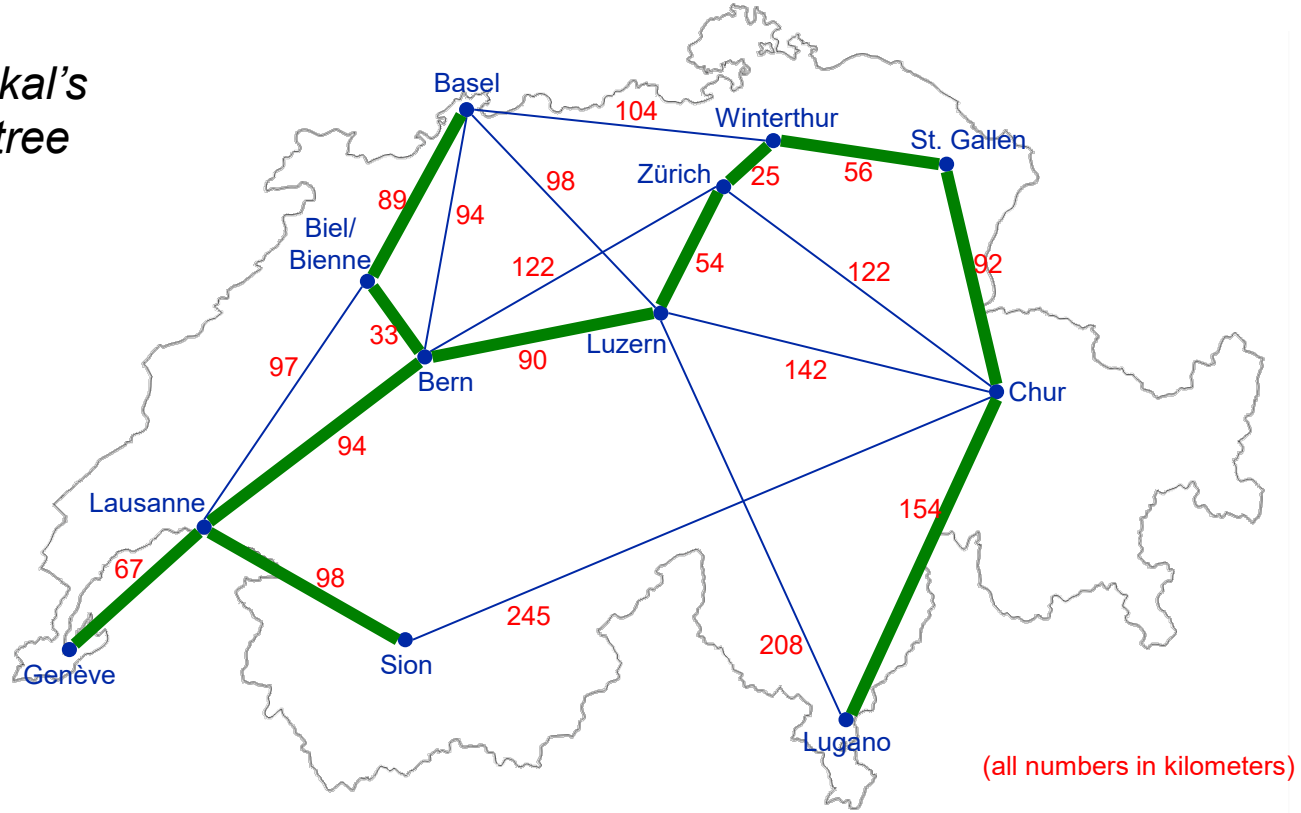

Dijkstra's Algorithm: Introduction

- «Dijkstra = BFS with looking back»
- Greedy: always does what currently looks best
- Finds shortest paths from a given start node to all other nodes (SSSP)

- Idea: Grows a set of finalized nodes for which the closest distance to the starting node is definitively known. Subsequently add new nodes reachable from the finalized set and than relax all edges leading to nodes from the finalized set of nodes.

- Uses a minimum priority queue (= min heap) using their distances as keys with the following two additional operations:
  - decreaseKey(u, d): change key (distance) of node u to new value d (reorder PQ accordingly)
  - extractMin(): removes node with minimal key (distance) from the queue and returns its value

Relaxation

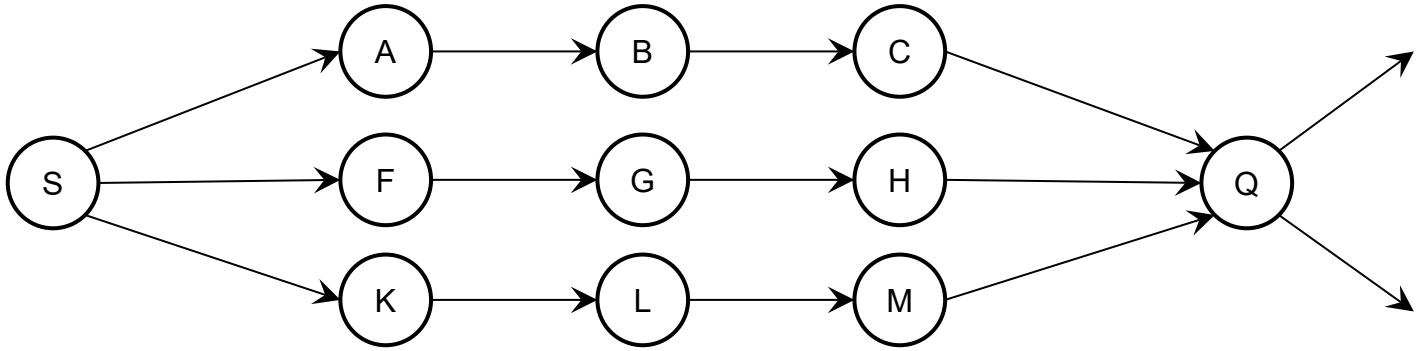

Dijkstra's Algorithm: Example

Apply Dijkstra's algorithm to find the shortest paths from node S to all other nodes:

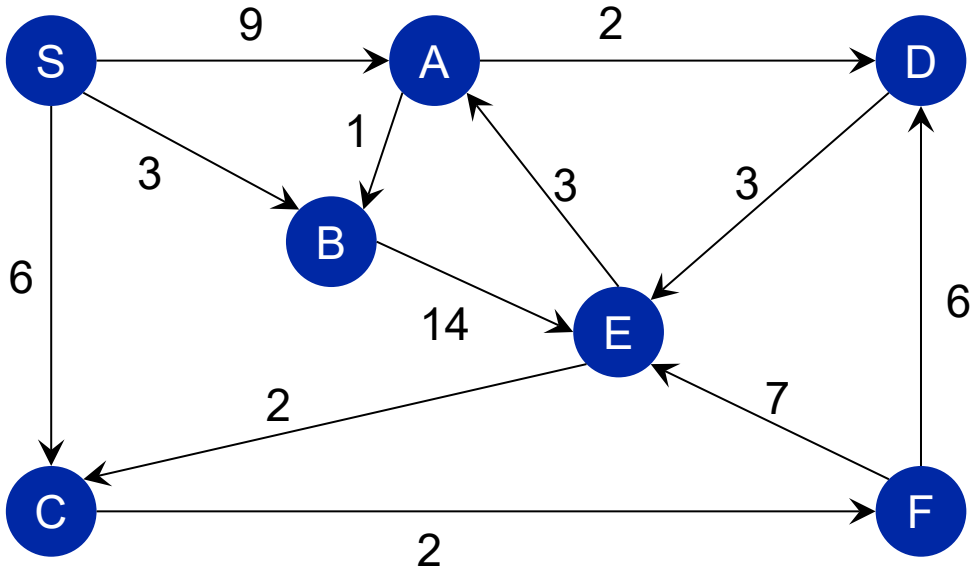

Dijkstra's Algorithm: Another Example

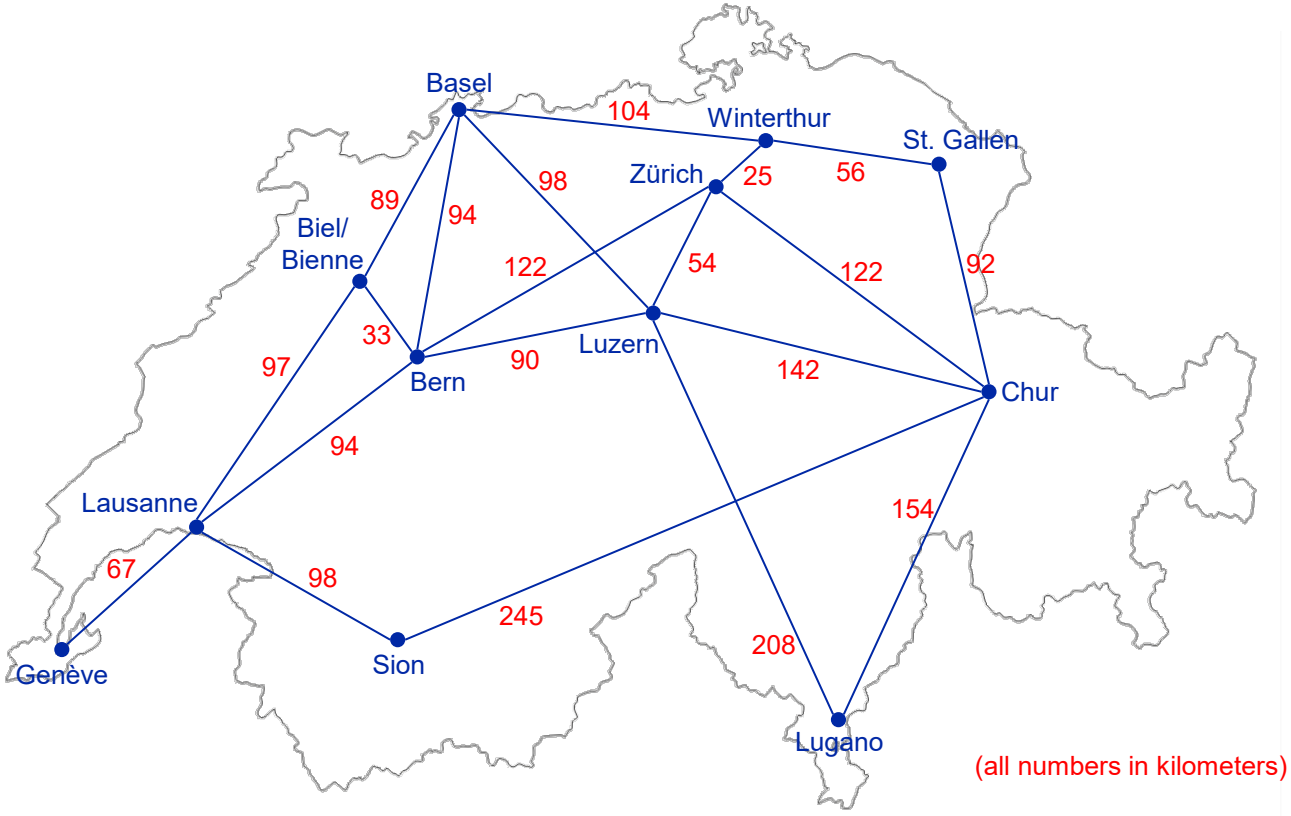

Dijkstra's Algorithm: Implementation

edge weights

```
Algo: Dijkstra(G,w,s)
foreach v ∈ G.V do
    v.dist = ∞; v.pred = NIL;
s.dist = 0;
InitMinPQ(PQ,G.V);
while PQ ≠ ∅ do
    v = ExtractMin(PQ);
    foreach u ∈ v.adj do
        Relax(u,v,w);
        DecreaseKey(PQ,u,u.dist)
```

as long as the priority queue is not empty

put all nodes into a minimum priority queue (with their distances as keys); priority queue will serve as a kind of to-do list (for all nodes which are still in the queue, the definite shortest path has not yet been found)

Is There Always a Shortest Path and Does Dijkstra Always Work?

Dijkstra's algorithm is not applicable (will fail) for some particular graphs with negative edge weights.

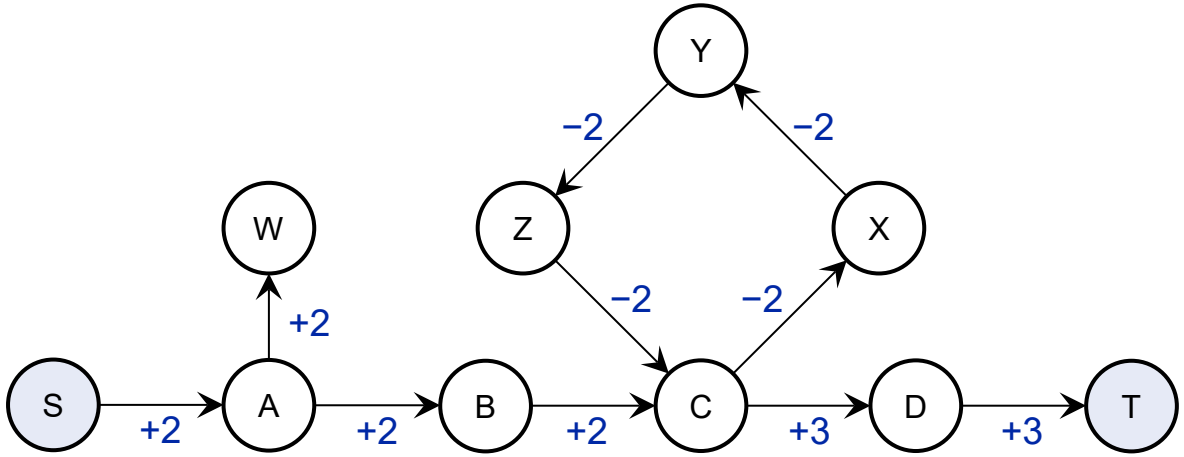

- The shortest path from S to A is 2.
- The shortest path from S to B is 4.
- The shortest path from S to W is 4.
- The shortest path from S to C is $-\infty$.
- The shortest path from S to all other nodes is $-\infty$.

– Why is this the case? What will happen if Dijkstra's algorithm is applied in this case?

– How can the shortest path from a single source be found in a graph with negative edge weights?

Bellman-Ford Algorithm: Implementation

Algo: BellmanFord(G,w,s)

```
foreach v ∈ G.V do
    v.dist = ∞; v.pred = NIL;
for i = 1 to |G.V|-1 do
    foreach (u, v) ∈ G.E do
        Relax(u,v,w);
foreach (u, v) ∈ G.E do
    if v.dist > u.dist + w(u, v) then
        return FALSE;
return TRUE;
```

Wrap-Up

- Summary
- Feedback
- Outlook
- Questions

Wrap-Up

– Summary

Questions?

Retrospective

https://www.klicker.uzh.ch/algodat

- Please take 5 minutes to answer these questions:
  - What were the most important points / insights of today's exercise session for you?
  - What things remained unclear or were confusing? About what do you want to know more or needs clarification?
  - What you wanted to say anyway…

Thank you for your attention.